

ECE 4300 | Computer Architecture

Image Processing Performance | FPGA vs CPU

Jessalin Jiangkhov, Rick Unite, Izaiah Adams

Why we chose this?

- Combines current learning of Computer Architecture with previous experience of digital design
- Hardware and Software focused
- Easily doable
 - Components already in possession
- Real Life Relevance
 - FPGAs used for parallelism and low latency
 - CPUs used for flexibility, easy to program, not as efficient for data-parallel tasks
 - Understanding trade-off prepares for modern computing challenges in AI hardware, embedded systems and accelerators



What is Image Processing?

Also known as Digital Image Processing, it is a process of analyzing and manipulating images digitally.

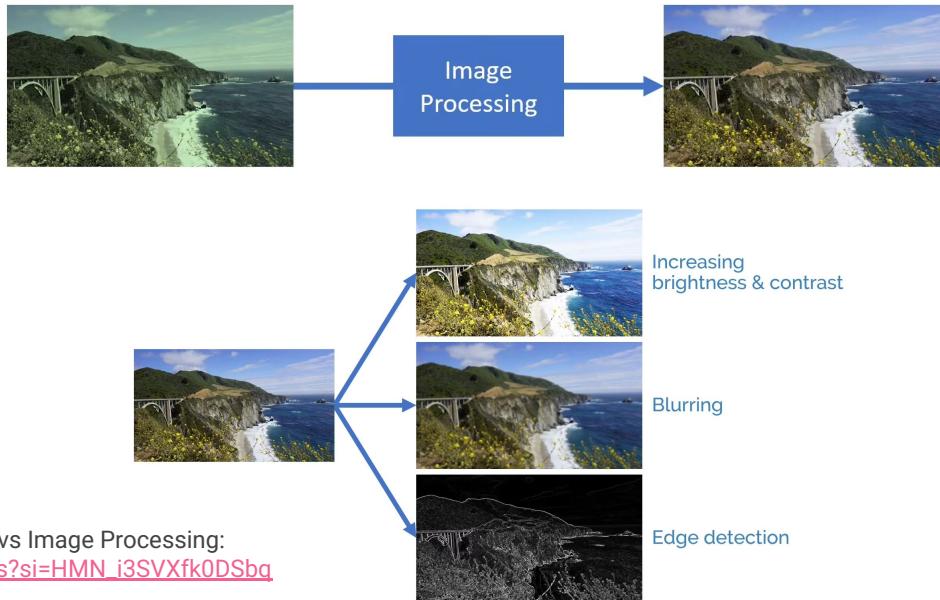


Image from Computer Vision vs Image Processing:
https://youtu.be/9-8Js62wzOs?si=HMN_i3SVXfk0DSbg

Objective

Testing and comparing the speed, power consumption, resource utilization, and output quality of an FPGA producing an image vs a CPU producing an image

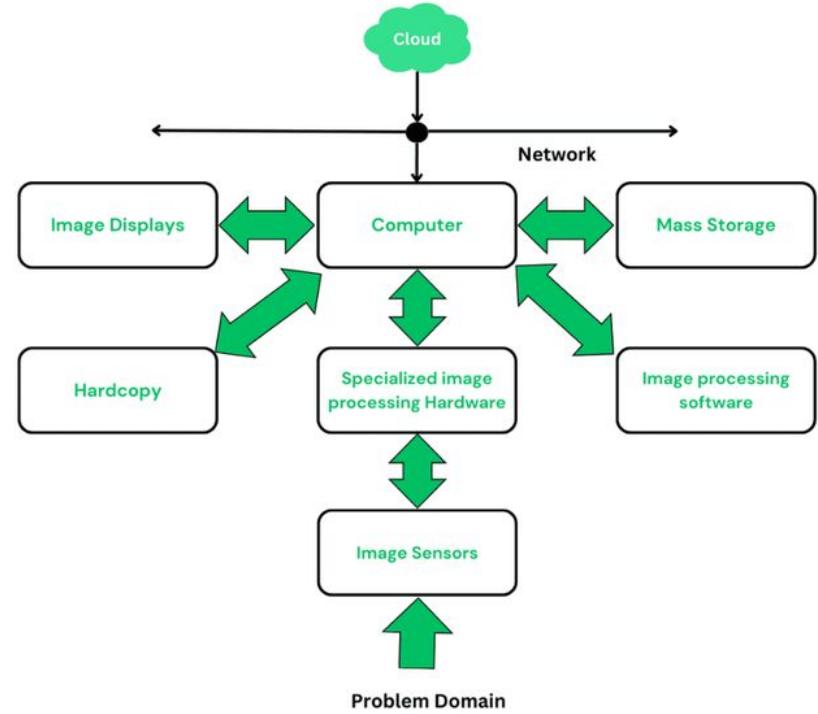


Image from geeksforgeeks:

<https://www.geeksforgeeks.org/electronics-engineering/what-is-image-processing/>

Hardware Background | FPGA

Field-programmable gate arrays (FPGA)

- Integrated circuits
- Can be reconfigured to users' desires
- Contain programmable logic blocks
- Flexible usage
- Hardware-timed speed and reliability
- Parallelism
- Found in:
 - Data Centers, aerospace & defense systems, medical devices, etc

Hardware Background | CPU

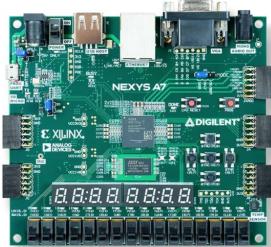
Central Processing Unit (CPU)

- “Brain” of a computer
- Executes instructions
- Parallel Processing
 - Modern CPUs utilize multicore processing to execute tasks simultaneously
- Clock Signal used for operation synchronization
- Multithreading
 - A single core can handle multiple sequences of instructions (multitasking)
- Found in many common devices
 - Computers, smartphones, smart TVs, etc

Hardware Specs

FPGA - NEXYS A7 100T

- 128MiB DDR2
- Serial Flash
- 4860 Kbits of fast block RAM
- Six clock management tiles
- Internal clock speed >450MHz
- 12-Bit VGA Output



CPU - Ryzen 9 7600x

- Zen 4 Processor Architecture
- 6 CPU Cores
- 12 Threads
- Max Boost Clock of 5.3GHz
- Base clock of 4.7GHz
- Graphics Model: AMD Radeon Graphics
- Graphics Core Count: 2
- Graphics Freq.: 2200MHz



Other Necessary Components

- Monitor capable of VGA (Video Graphics Array) Display for the FPGA output
- Working computer capable of using the desired CPU
- OV7670 camera module for FPGA Camera capabilities
- Vivado 2025.1 will be used to program FPGA



OV7670 Camera
Module

Testing specifics: Image processing algorithms

- Simple, compute-heavy kernels - ideal for both CPU and FPGA
 - Grayscale conversion
 - RGB → grayscale using weighted sum (simple, tests memory throughput)
 - Image convolution
 - Apply filters like blur, sharpen, etc. (great for pipeline parallelism)
 - Sobel edge detection
 - Gradient computation (x, y directions) (common benchmark for edge detection)
 - Gaussian blur
 - Smooth image; uses 2D convolution (compute-intensive, good for throughput test)
 - Thresholding / binarization
 - Convert image to black/white by pixel intensity (easy control and branching test)
 - Matrix multiplication (for comparison)
 - Core operation for CNN-like workloads (useful for performance scaling)

Testing specifics: Performance metrics

- Measure quantitatively comparable results for both platforms:
 - Execution Time / Latency
 - Wall-clock timing (`time.perf_counter()` (CPU) / on-chip timer (FPGA))
 - Throughput (FPS)
 - $1 / \text{frame time}$ (Manual calculation)
 - Resource Utilization
 - LUTs, FFs, BRAM, DSP usage (Vivado/Quartus synthesis report)
 - Power consumption
 - Estimated power (Xilinx Power Analyzer / Quartus Power tool)
 - Energy per frame
 - Power \times time (Derived metric)
 - Clock frequency
 - Achieved timing (MHz) (Synthesis report)

Testing specifics: Datasets

- Open source dataset:
 - Kodak dataset (24 uncompressed images)
 - Berkeley Segmentation Dataset (BSDS500)
 - COCO / ImageNet samples (downscaled)
 - Or capture our own test images
 - (MUST USE SAME IMAGES FOR BOTH CPU AND FPGA)

Testing Process

- 1.) Select algorithm
- 2.) Implement on CPU using OpenCV
- 3.) Implement on FPGA using HLS (C code -> hardware)
- 4.) Run both on same dataset (same image size, e.g 512 x 512 or 1080p)
- 5.) Record results (time, fps, LUTs, power)
- 6.) Plot comparison graphs: latency vs throughput, resource vs performance

Testing Process | FPGA

1. Testing the FPGAs capabilities using an algorithm to produce an image to a VGA Monitor
 - a. Will reference this tutorial: <https://www.instructables.com/Image-from-FPGA-to-VGA/>
2. Testing the FPGAs capabilities of outputting the live feed of a camera
 - a. Will reference this GitHub: <https://github.com/amsacks/OV7670-camera>
 - i. User's Demo Video: <https://youtu.be/PbZsmB2INCU?si=O62HTviTFL-5ToOY>
3. Record Data and Outputs
 - a. Take pictures of Image / Video Outputs
 - b. Note device temperature

Testing Process | CPU

TBD

Current Status

- Most components obtained
 - Further research was necessary
 - Must order camera for FPGA
- Testing will be conducted soon
 - Availability Issues
- Benchmarking and other tests will be conducted on separate devices
 - Results will later be compared
- Recovered from a hardware failure

Data and Results

- Typical Trend (Hypothetical)
 - FPGA: 3–10× faster for small kernels when pipelined properly.
 - CPU: Easier to implement, but slower and more power-hungry.
 - FPGA: Requires more setup but demonstrates parallelism.

Issues and Other Possible Variables

- **Hardware-related issues**
 - FPGA resource limitations: not enough LUTs, BRAM, or DSP slices for large images or complex filters
 - Power and Heat Constraints: Running the FPGA at high frequency or full utilization may cause overheating
 - Hardware interface or Driver Issues: Data transfer between CPU and FPGA can be buggy or slow
- **Software and Toolchain problems**
 - Toolchain complexity: Steep learning curves
 - Simulation vs Real Hardware Mismatch: Design works in simulation but fails on the actual board
 - Timing and Optimization: FPGA synthesis succeeds but achieved frequency is lower than expected
 - CPU Version Optimization Bias: CPU version might run slower or faster depending on how we implement it
- **Benchmarking & Measurement Pitfalls**
 - Unfair Comparison: (different resolutions, data movement overhead not included)
 - Inaccurate Timing: Measuring performance incorrectly (I/O or OS overhead)
 - Power Estimation Uncertainty: Estimating FPGA power = hard without a physical meter
- **Algorithmic and Data Issues**
 - Large Image Sizes: might exceed FPGA memory
 - Fixed-Point Precision Errors: converting floating point to fixed point in FPGA = output accuracy may degrade

